FlexPoint Forecast Model — v2

Full Speaker Script

Wednesday, February 25, 2026

Liam Harvey & Ajer Sher | Gallus Insights

Audience: Augie, Mariana, Israel (FlexPoint)

Target: 15 min presentation + 15 min discussion

General Notes Before the Call

Timing

Slides 1-8: ~10 minutes of presentation. Slide 9-10: ~5 min framing + open discussion.

Don't rush the backtest slide (6) or the lock diagnostic slide (4) — those are where Augie will have the most questions.

If running short, skip Slide 5 (Feature Importance) — Ajer can summarize it in one sentence during the lock slide.

Tone

Augie is sharp and detail-oriented. He'll pull out his Excel blackboard if something doesn't click. That's a good thing — lean into it.

Mariana and Israel are more operational. Keep it intuitive for them, but don't dumb it down. Augie sets the bar.

When you say 'you asked for this' or 'your insight' — mean it. Augie responds well to genuine credit, not flattery.

Numbers to Have in Your Head

15,272 total loans in the dataset (Jan 2023 - Jan 2026)

5,661 funded loans (37% pull-through)

13 pipeline stages, Opened through Funded

v1 MAPE: 6.4% at mid-month, 18/24 within 10%

v2 MAPE: 4.8% at mid-month, 23/24 within 10%

Lock coverage: 0% at Opened, 33% at Approved, 100% past Condition Review

99.3% of funded loans are locked. 0.6% of unlocked loans fund.

Seasonality: summer pull-through ~41%, winter ~36%. Volume +25% higher in summer. Days-to-fund is flat.

Slide 1 — Title

~15 seconds. Ajer opens.

[AJER]

"Thanks everyone for joining. Liam and I are going to walk you through the v2 model update today."

"Quick agenda — we'll show how we built on the feedback from last week, what the results look like across 24 months, and then we have a couple of open questions where we want your input."

Click to next slide.

Slide 2 — What We Changed

~45 seconds. Liam presents. This is setup — move through it briskly.

[LIAM]

"So from our last call, Augie, you laid out four things you wanted us to work on. We did all four. Let me walk through them quickly."

"First — separate product and purpose. Last time, product type and loan purpose were buried inside the stage probability. You said you wanted to see what each one contributes on its own. Done. They're now independent features."

"Second — rate lock. You spent the last ten minutes of the call explaining how the lock date and lock expiration date could be really powerful. We built five new features around that. We'll get into the results in a few slides."

"Third — the day-zero projection. You asked us to test what happens if we stand on the last day of the prior month and project the entire next month forward. So for December, we stand on November 30th. For November, we stand on October 31st. Each month uses the most current pipeline available. No stale data. That's now built in."

"And fourth — the rolling window. Instead of always training on 2024 data, we tested a rolling 12-month window so the model adapts when conditions change. This turned out to be a big deal — we'll show you why."

"So those are your four requests. Let's look at what we found."

Slide 3 — Product & Purpose

~90 seconds. Ajer presents. This is the 'Augie was right' slide.

[AJER]

"Augie, you were right. Adding product and purpose makes a real difference in the probabilities."

"So what you're looking at here is funding probability by pipeline stage — the gray bars are the baseline, and then the colored bars break it out by product and purpose."

"Focus on the Approved stage in the middle — that's where most of our active loans sit. The baseline is about 52 percent. But when you break it out, a nonconforming purchase loan funds 60 percent of the time. FHA purchase is even higher — 64 percent. And a refinance cashout is down at 47."

"And at Submitted on the left, the spread is even wider. Nonconforming loans are nearly double the baseline — 39 versus 21 percent."

"Now, one thing to flag honestly — when we separate them out as independent features in the model, their standalone importance is small. Most of that signal was already captured by the pipeline stage. But the key thing is they're visible now. You can point to them and say, this is what product contributes, this is what purpose contributes. Which is what you asked for."

If asked: "How did you calculate these?"

> *"We built stratified transition tables from all 15,000-plus loans. For each combination of stage, product, and purpose, we computed the historical funding rate. Where cell sizes were too small — under 20 loans — we smoothed it back to the stage-only probability so we're not overfitting to a handful of loans."*

If asked: "What about refinance versus purchase broadly?"

> *"Purchase loans fund at about 44 percent overall. Refinance is 25 percent. Refi cashout is in between at 35. Purpose is the single most impactful stratifier after pipeline stage."*

Slide 4 — Rate Lock Diagnostic

~90 seconds. Liam presents. This is the nuanced slide — take your time.

[LIAM]

"This is where it gets interesting. This chart shows what percentage of active loans have a rate lock at each pipeline stage."

"On the left — Opened, Application — basically zero percent are locked. That makes sense. Nobody's locking a rate on a loan that just came in the door."

"On the right — Condition Review, CTC, Docs — a hundred percent are locked. By that point, every loan has a locked rate."

"The magic is in the middle. At Approved, about a third of loans are locked and two-thirds aren't. That split is genuinely useful — it separates loans that are actively moving forward from ones that are sitting there."

"Now here's the honest part. Ninety-nine percent of loans that eventually fund have a rate lock. And less than one percent of unlocked loans end up funding. So the lock is almost a perfect predictor — but that's partly because it's telling us the same thing the pipeline stage already tells us. A locked loan is just further along."

"The lock features improved our mid-month accuracy from 6.5 percent down to 4.8. That's a meaningful jump. But 68 percent of the model's weight ends up on lock timing. We'll come back to what that means and ask for your take on it."

If asked: "Is this the same thing I described with Loan A and Loan B?"

> *"Exactly. The number-one feature the model learned is lock expiry versus month-end — which is precisely your example. A loan whose lock is about to expire is under pressure to close. A loan with plenty of lock time left is on track but not urgent. The model picked up on that pattern directly."*

Slide 5 — Feature Importance

~60 seconds. Ajer presents. If running short, this can be summarized in 20 seconds and folded into Slide 4.

[AJER]

"So here's what the model actually relies on, ranked from top to bottom."

"Lock expiry versus month-end — that one feature carries 68 percent of the weight. That's the signal Augie described. How close is the lock expiration to the end of the month? If it's expiring soon and the loan hasn't funded, something has to give."

"Days at the current stage is second at about 7 percent. A loan that's been sitting at Approved for 45 days is a very different animal than one that just arrived five days ago."

"Pipeline stage probability — which was 60 percent of the weight in v1 — dropped all the way down to 6 percent. The lock features basically absorbed that signal."

"And then credit score, loan amount, interest rate — they're all in the low single digits. Useful at the margins, but the heavy lifting is done by lock timing and stage progression."

If asked: "Where are product and purpose on this chart?"

> *"Same story as before — their signal is already captured upstream in the stage probability. They matter, they're just embedded. If you ran the model without them, the stage probability would shift to compensate. They're visible as independent features now, but their standalone weight is small."*

Slide 6 — Backtest Results

~90 seconds. Liam presents. This is the proof slide. Don't rush.

[LIAM]

"This is the proof. We backtested the model across 24 months — January 2024 through December 2025."

"Three numbers at the top. The first — 4.8 percent average error at mid-month. That means if FlexPoint funds a hundred million dollars in a month, we're typically off by about 5 million. That beats the 10 percent target by a wide margin."

"The second — 23 out of 24 months land within 10 percent of actual. In v1, it was 18 out of 24. So we went from three-quarters of months hitting the target to basically all of them."

"And the third — that's a 25 percent reduction in error compared to v1. We went from 6.4 down to 4.8."

"The chart below shows it visually. The black line is what actually funded each month. The orange dashed line is our v2 projection. You can see they track tight — especially in the second half of 2025 where volumes ramped up significantly and the v1 model started struggling."

"The one month still outside the target zone is November 2025. Let's look at that next."

If asked: "What about early in the month — day 1, day 8?"

> *"Day 1 is about 16 percent error — you're projecting almost the entire month, so it's naturally harder. By day 8, we're under 10 percent. By day 15, 4.8. And by day 22, we're at 4 percent. The model gets sharply better as actuals come in."*

If asked: "Which model won?"

> *"Gradient boosting won on every metric — lowest error, best calibration, highest AUC. We also tested logistic regression, random forest, and LightGBM. Logistic was close on calibration but couldn't handle the nonlinear time-at-stage effects."*

Slide 7 — Nov/Dec Deep Dive

~60 seconds. Ajer presents.

[AJER]

"November and December were our two worst months in v1. This slide shows how they improved."

"November — in v1, we over-projected by 28 percent. With v2 on a fixed training window, it came down to 18. And with the rolling window, it came down to about 13. Still above 10 percent, but we cut the error by more than half."

"December is the better story. V1 was off by 15 and a half percent. V2 fixed brought it to 8. Rolling window brought it to 6. That's inside the 10 percent target."

"The key insight here — the rolling window is doing the heavy lifting. When the model trains on the most recent 12 months, it's seen the higher volumes from mid-2025. The fixed window was still calibrated to 2024 levels, where monthly volumes were 60 to 120 million. By late 2025, volumes were pushing 150 to 175."

"And this improvement has nothing to do with the lock question. The rolling window is a clean win. We'd recommend it regardless of what you decide about locks."

If asked: "Why is November still above 10 percent?"

> *"November actual was about 123 million against a pipeline that looked like 140 million at mid-month. Something caused an unusually high drop-off that month — possibly year-end slowdown starting early, holiday closures, or loans being held for January funding. We'd love your input on whether there's a known pattern there."*

Slide 8 — Snapshot Day Accuracy

~60 seconds. Liam presents.

[LIAM]

"This is the new slide you asked for. It shows how accurate the model is depending on when in the month you run it."

"On the left — Day 0. That's the test you wanted. We stand on the last day of the prior month and project the entire next month. Nothing has funded yet. It's all forward-looking. At that point, we're at about 17 percent error. That's a directional read — good enough to know roughly where the month is heading, but not precise."

"By Day 8 — one week of actuals are in. Now we're under 10 percent."

"Day 15 — half the month is actual, half is projected. That's where we hit 4.8 percent, which is our headline number."

"And Day 22 — most of the month is known at that point. We're at 4 percent."

"The important thing structurally — every single month is projected from its prior month-end. December from November 30th. November from October 31st. Always the freshest pipeline available. That's the consistent framework you laid out."

If asked: "Why is the rolling window worse at Day 0?"

> *"Good catch. The rolling window is optimized for mid-month accuracy. At Day 0, it's actually more aggressive with its probabilities because it's seen higher recent volumes and assigns higher funding rates. The fixed window is more conservative early on, which happens to overshoot less at Day 0. But by Day 8, the rolling window catches up and wins from there forward."*

Slide 9 — Two Paths Forward

~2 minutes. Ajer frames, Liam asks the question. This is where the discussion starts.

[AJER]

"So we want to lay out two options and get your input."

"On the left — Version A. The conservative version. No lock features. 6.5 percent mid-month error. It's clean, it's defensible, there's no question about whether the model is just learning something it already knows from the pipeline stage. Product and purpose are independent. Rolling window is included. This is the safe choice."

"On the right — Version B. With lock features. 4.8 percent mid-month error. Significantly better accuracy. But 68 percent of the model's weight is sitting on lock timing. And as we showed, lock status is closely tied to pipeline stage — it might be telling us what we already know, just in a different way."

[LIAM]

"So here's the question for you, Augie. You understand the capital markets side better than we do. In your experience — is the rate lock a genuinely independent signal? Does locking cause a loan to get priority processing, or does it get a different kind of attention? Or does a loan just get locked along the way because it's already progressing, and the lock is basically a byproduct?"

"The one place we know it adds real information is at Approved, where about a third are locked and two-thirds aren't. That split separates movers from starters. But past Approved, everyone's locked, so it's not telling us much new."

"Either way we're comfortable. Both versions use the rolling window and the separated product and purpose features. We just want to make sure we're building on a foundation you trust."

[PAUSE — Let Augie respond. This is the most important moment in the call.]

If asked: "Can we use it? / Go with Version B."

"Great. We'll move forward with Version B. We'll document the lock dependency clearly so anyone reviewing the model understands where the weight is concentrated."

If asked: "I'm not comfortable with it. / Go with Version A."

"Understood. Version A is still a meaningful improvement over v1 — the rolling window alone cuts November and December errors dramatically. And the product-purpose separation gives you the auditability you wanted."

If asked: "I'm not sure. / Can we test both?"

"Absolutely. We can deploy Version A now as the production model and run Version B in parallel as a shadow. That way you can see how they compare on live data for a month or two before committing."

Slide 10 — Next Steps

~60 seconds + open discussion. Liam presents.

[LIAM]

"Four things going forward."

"First — your call on lock features determines which version we ship. Rolling window is in either way, so that improvement is locked in, no pun intended."

"Second — seasonality. You flagged this last week. We ran the numbers. Summer pull-through is about 41 percent versus winter at 36. Volume is 25 percent higher in summer months. But here's the interesting part — the actual days to fund is flat year-round. It's not that loans close faster in summer. It's that more of them close. The conversion rate shifts, not the speed. That's a natural next feature to add to the model."

"Third — we have the full error-by-month chart showing all three versions side by side if you want to dig into any specific month. Happy to walk through it."

"Fourth — integration. The scoring module already produces the same output format as what's in ThoughtSpot today. Connecting it to the production database is about a one to two week effort. We don't need to start that yet, but the path is clear."

[AJER]

"And that's it from us. Happy to go deeper on anything or take questions."

[Open discussion begins. Let Augie lead.]

Backup Q&A — If It Comes Up

These are questions that might come up during open discussion. Both Liam and Ajer should be ready.

If asked: "When can this go live?"

"The model itself is code-complete. The integration step is connecting it to the production database and setting up a daily scoring job. That's about one to two weeks depending on database access and coordination with the IT side."

If asked: "Can you predict further out — like next month from this month?"

"Same architecture. We'd just extend the time horizon. Right now we predict month-end from any point in the current month. Predicting next month from today would use the same pipeline snapshot, same scoring logic — we'd just need to build transition tables calibrated for a 30 to 60 day horizon instead of remaining days in this month."

If asked: "What about the prior model — the one with 99 percent accuracy?"

"That was a classification model — it asked 'can I tell apart funded loans from canceled ones after the fact?' Which is a much easier question. It used features like time from Approved to CTC, which only exists for loans that already reached CTC. So it was using the answer to predict the answer. We rebuilt from scratch with a forecasting mindset — given what I know today, what will fund by month-end?"

If asked: "What if lock features start degrading over time?"

"That's a real concern. If the relationship between lock and funding changes — say locks get extended more frequently, or lock policies change — the model would need to adapt. The rolling window helps with that naturally since it's always training on the most recent data. But we'd monitor feature importance monthly to catch any drift."

If asked: "Can you show individual loan scores?"

"Yes. The model produces a loan-level detail table — every active loan with its probability and expected dollar contribution. We can drill into any individual loan and say, here's why the model thinks this one has an 80 percent chance or a 20 percent chance."

If asked: "What about the retail-wholesale split?"

"The portfolio is overwhelmingly wholesale — about 96 percent. Retail only contributes about 5 to 6 million a month. Wholesale funds at 45 percent, retail at 11 percent. The model handles both but wholesale drives the numbers."

If asked: "Is there a year-end effect we should worry about?"

"That's part of what we want to ask you. November and December consistently show higher errors. It could be holiday closures slowing down late-stage processing, lender capacity constraints at year-end, or loans being intentionally held for January. If there's a known pipeline freeze or slowdown, that would explain the pattern and we could build it in."